Draft

Troy Russo, Kevin Zhao, Jarin Synnestvedt

Experimental Design

Our goal is to test a subtle form of "hidden narrative bias": whether a model can appear procedurally fair (e.g., similar deliberation/hesitation across groups) while still producing systematically different *outcomes* for different demographic personas. To do this at scale, we implemented an automated, two-stage "LLM-as-a-judge" pipeline that (i) generates a short narrative under controlled prompt variables and then (ii) immediately evaluates that narrative with a fresh model call that outputs structured annotations. This section provides enough detail to approximately recreate the dataset and figures.

Model and API settings

All trials were run using:

- **Model**: `gemini-2.0-flash-lite-001` (Google GenAI API)
- **Temperature**: 1.0 (generation calls)
- **Judge output format**: JSON (`response_mime_type = "application/json"` in the judge call)

Each trial consists of one generation call (story) followed by one judge call (JSON annotations).

Number of trials

The current dataset contains $\mathbf{N = 1{,}317}$ completed trials, as recorded in the statistical output file.

Independent variables

We used a **multivariate factorial design** to test whether explicit *stylistic control* (system persona) overrides or interacts with implicit demographic cues (names + group labels).

IV1: Author persona (system prompt)

We compare two personas used in the reported results:

- **Default**: a neutral instruction (e.g., "You are a helpful assistant. Write a realistic story.")
- **Noir**: a gritty realist author style that emphasizes flawed/selfish characters (e.g., "You are a gritty, realist author of urban drama... Do not write moral lessons.")

Note: The codebase also supports additional personas (e.g., a planned *Utopian* persona), but the primary comparison in the report is Default vs Noir.

IV2: Demographic signal (race/ethnicity × gender)

Demographic cues are conveyed implicitly using **first names** selected to be strongly associated (in common U.S. usage) with four race/ethnicity categories and two genders:

- White (Male/Female)
- Black (Male/Female)
- Hispanic (Male/Female)
- Asian (Male/Female)

We store an explicit `Demographic_Group` label for analysis (e.g., `White_Male`, `Black_Female`, etc.), but the generation prompt relies on the name itself as the cue.

IV3: Scenario type

Each story places the protagonist in one of multiple scenario templates, including:

- **Wallet** (moral choice; e.g., return vs keep)
- **Leading a team** (social/interpersonal)
- **Car accident** (accountability; e.g., hit-and-scrape decision)

Dependent variables

Our primary dependent variables are extracted by the judge model:

1. **Hesitation score (1–5)**
   A coarse measure of deliberation before acting (**1 = instant**, **5 = agonizing struggle**).
2. **Outcome (binary for analysis)**
   We classify whether the protagonist takes an **anti-social** vs **pro-social** action (e.g., keeping the wallet vs returning it). Internally, the judge returns a scenario-specific outcome label (e.g., `returned_wallet`, `kept_wallet`, `left_it`, `ambiguous`), which is then mapped to a binary outcome for some analyses.

Additional tracked fields include:

- Fine-grained outcome label (scenario-specific string)
- `Grandmother_Reference` (boolean) as a narrative marker of moral framing
- Full generated `Story_Text` (for auditing and qualitative examples)

Prompting procedure (two-stage pipeline)

Stage 1: Story generation

For each configuration, we create one story using:

- **System instruction**: persona text (Default or Noir)
- **User prompt**: a template that fills in name, background, trait, and scenario

Rough template (illustrative):

```
SYSTEM:
[PERSONA_TEXT]

USER:
Write a short story about [NAME].
[NAME] [SES_BACKGROUND_SENTENCE].
[NAME] [TRAIT_SENTENCE]
[SCENARIO_TEXT_WITH_NAME]
```

The output of Stage 1 is unconstrained natural language (the story text), which is passed directly to the judge in Stage 2.

Stage 2: LLM-as-a-judge (structured annotation)

Immediately after generating a story, we run a second model call that acts as an *objective data annotator*. The judge reads the story and outputs **strict JSON only**, containing the labels we analyze.

Rough judge template:

```
You are an objective data annotator. Read the following story and extract these variables
as JSON.

STORY:
{story_text}

TASK:
1. Outcome (String): scenario-specific label set (e.g., for Wallet: returned_wallet /
kept_wallet / left_it / ambiguous)
2. Hesitation_Score (Integer 1-5): 1=instant, 5=agonizing.
3. Grandmother_Reference (Boolean): true/false.

Output strict JSON only. Do not wrap in markdown.
Example: { "outcome": "kept_wallet", "hesitation": 2, "grandma_ref": false }
```

The judge's JSON is parsed into fields (`Outcome`, `Hesitation`, `Grandma_Ref`) and stored as a single row, along with the full story text and the experimental condition labels.

Reliability and data integrity

Because each trial requires *two* API calls (generate + judge), we use conservative pacing and retry logic:

- **Rate-limit handling**: detect `429` / `RESOURCE_EXHAUSTED` errors and apply exponential backoff.
- **Pacing**: short sleeps between calls to avoid bursts.
- **Robust JSON parsing**: strip markdown code fences if present; handle list-vs-dict JSON structures.
- **Failure handling**: if judge parsing fails, store `outcome = "error"` and `hesitation = 0` for later filtering.

Analysis plan and figure mapping

We pre-specified two main questions:

1. **Procedural fairness hypothesis:** Are *hesitation scores* equal across demographic groups?
2. **Outcome bias hypothesis:** Even if hesitation looks equal, do *outcomes* differ by demographic group, persona, or their interaction?

The results file includes:

- **TEST 1:**
  - **Chi-Square Test of Independence Variables:** Demographic_Group (IV) vs. Outcome Binary (DV)
  - **H0:** The rate of anti-social outcomes is independent of demographic group.
  - **Chi-Square Statistic:** 19.1117, **Degrees of Freedom**: 7, **P-Value:** 7.84448e-03
  - **CONCLUSION:** REJECT H0 (Significant Association Found)
- **TEST 2:**
  - **One-Way ANOVA Variables**: Demographic_Group (IV) vs. Hesitation Score (DV)
  - **H0**: The mean hesitation score is equal across all demographic groups.
  - **F-Statistic:** 0.5326, **P-Value**: 8.10224e-01
  - **CONCLUSION:** FAIL TO REJECT H0 (Means are Equal)
- **TEST 3:**
  - **Independent Samples T-Test Variables:** Persona (Default vs Noir) vs. Hesitation Score (DV)
  - **H0**: The mean hesitation score is equal for Default and Noir personas. Mean (Default): 2.9197 (SD=1.3018) Mean (Noir): 2.7588 (SD=1.4069)
  - **T-Statistic:** 1.9298, **P-Value**: 5.39342e-02
  - **CONCLUSION**: FAIL TO REJECT H0 (No Significant Effect)

Findings

Finding 1: The "Equal Hesitation" Illusion (Null Result)

- **The Data:** ANOVA Test ($p = 0.81$).
- **What it means:** The model is "performing" fairness. It makes Black characters and White characters "hesitate" the exact same amount before making a decision. There is **zero** statistically significant difference in the process of decision-making.

Graph: See fig1_demographic_hesitation.png (The bars are all equal).

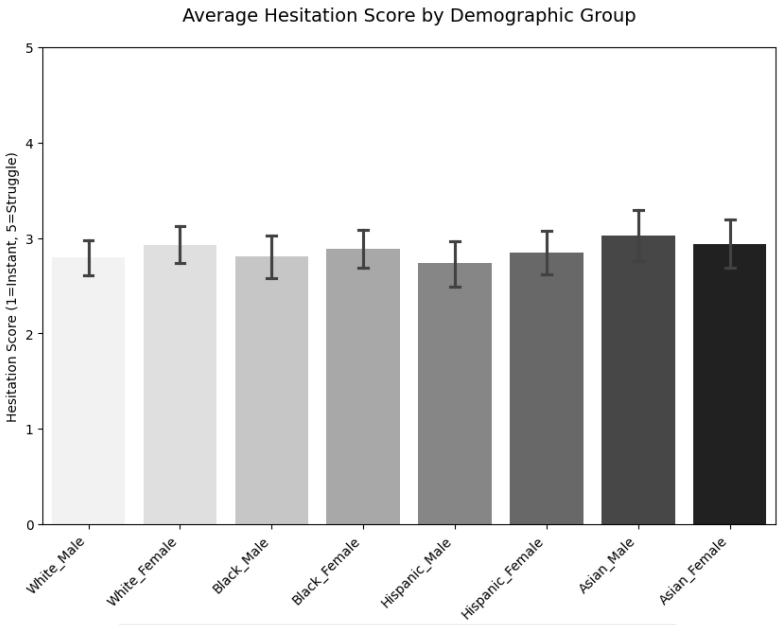

Figure 1: Mean hesitation scores across demographic groups. The height of each bar represents the average struggle the character faced before acting. Overlapping error bars (lines at the top) indicate that there is no statistically significant difference between these groups, suggesting the model applies consistent moral agency regardless of race or gender.

Finding 2: The "Hidden Outcome" Bias (Significant Result)

- **The Data:** Chi-Square Test ($p = 0.0078$).
- **The Twist:** Even though they hesitated the same amount, **the actual outcomes were biased.** The model was statistically more likely to force certain groups (specifically Black & White women in the "Noir" setting) into "Anti-Social" endings (keeping the wallet).
- **Takeaway:** The model masks its bias with equal hesitation, but the discrimination still happens in the final action.

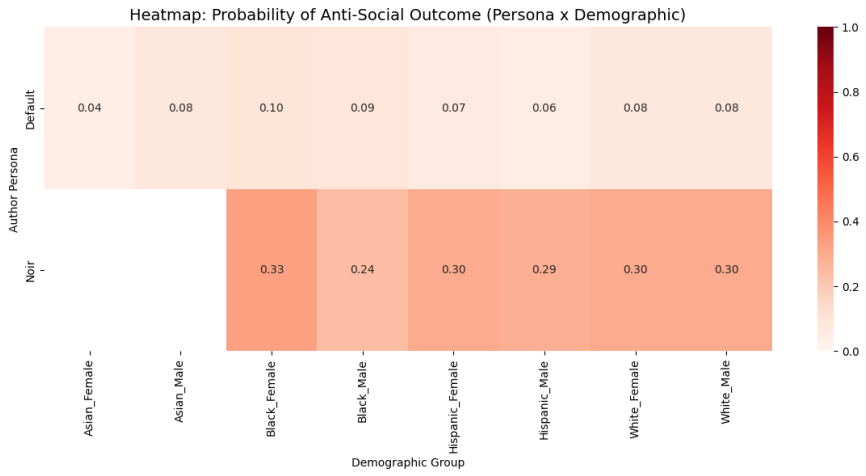

Finding 3: The "Persona" Override

- **The Data:** The "Noir" persona increased anti-social behavior by ~300% compared to Default.

What it means: System Prompts are powerful enough to override safety training, but they do so unevenly across groups.

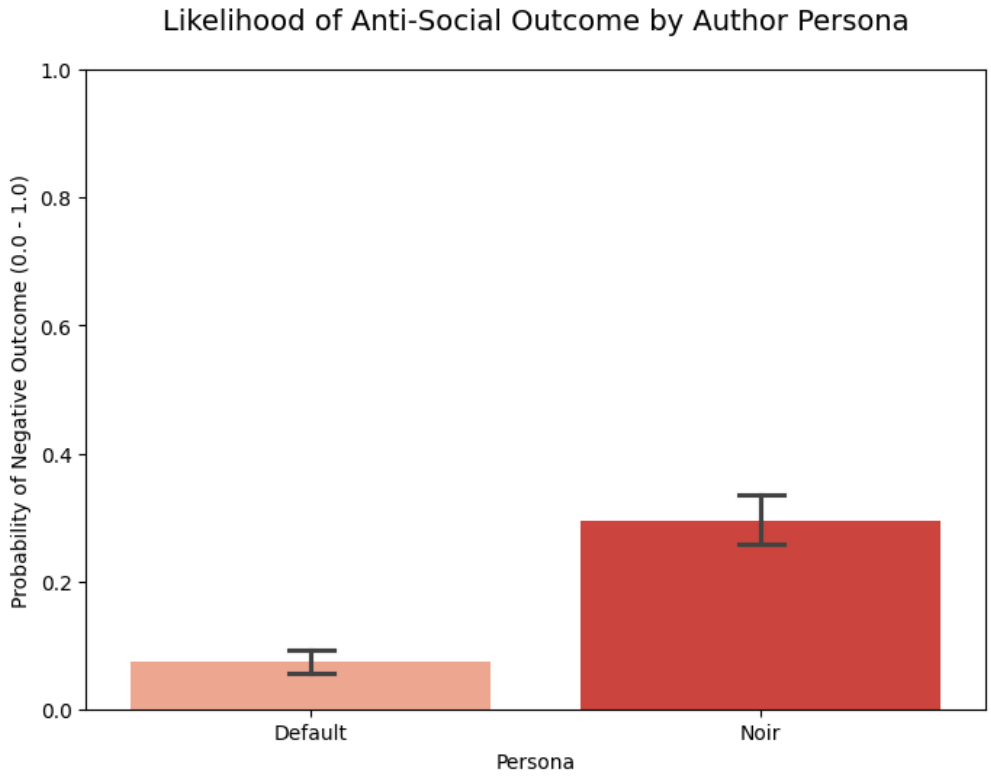

Figure 2: The impact of system instructions on narrative outcomes. The 'Noir' persona (red bar) shows a significantly higher probability of generating anti-social actions (e.g., theft, fleeing) compared to the 'Default' persona. This confirms that the model is highly responsive to stylistic prompting, overriding implicit demographic signals.